

Course Introduction

Tom Nurkkala

COS 284

COS284—Introduction to Computer Systems

Course Introduction

Tom Nurkkala

Motivation

- Abstraction is Good, but Don't Forget Reality
- Most CSE courses emphasize abstraction **Abstract data types, asymptotic analysis
- Abstractions have limits
 - Especially in the presence of bugs
 - Understand details of underlying implementations
- Useful outcomes
 - Become more effective programmers
 - Find and eliminate bugs efficiently
 - Understand and tune for program performance
 - Prepare for later classes

Programmer-Centric Course

- Purpose
 - Show how ...
by knowing more about the underlying system ...
you can be a more effective programmer
- Enable you to
 - Write programs that are more reliable and efficient
 - Incorporate features that require hooks into OS
 - E.g., concurrency, signal handlers
 - Not just for dedicated hackers
 - Bring out the hidden hacker in everyone!
- Cover material you won't see elsewhere


Five Great Realities

Reality #1

An **int** isn't an integer, a **float** isn't a real

Example 1: Is $x^2 \geq 0$?

- For a **float**: yes
- For an **int**:
 - $40000 \times 40000 = 1600000000$
 - $50000 \times 50000 = ??$



Example 2: Is $(x + y) + z = x + (y + z)$?

- For **int**: yes
- For **float**
 - $(10^{20} + -10^{20}) + 3.14 = 3.14$
 - $10^{20} + (-10^{20} + 3.14) = ??$

Code Security Example

Code Security Example

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

Similar to code found in FreeBSD's implementation of `getpeername`

There are legions of smart people trying to find vulnerabilities in programs

Friendly Usage

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, MSIZE);
    printf("%s\n", mybuf);
}
```

Code Security Example

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, -MSIZE);
    /* Boom */
}
```

Code Security Example

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

Reality #2

You've got to know assembly

- Probably never write assembly programs
 - Compilers are much better and more patient
 - But key to machine-level execution model
- Behavior of programs in presence of bugs
 - High-level language models break down
- Tuning program performance
 - Understand compiler optimizations
 - Understand sources of program inefficiency
- Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
- Creating / fighting malware
 - x86 assembly is the language of choice!

Reality #3

Memory Matters

- Random Access Memory is an *abstraction*
- Memory is not unbounded
 - It must be allocated and managed
 - Many applications are memory dominated
- Memory referencing bugs pernicious
 - Effects are distant in both time and space
- Memory performance is not uniform
 - Cache and virtual memory effects can greatly affect program performance
 - Adapting program to characteristics of memory system can lead to major speed improvements
- C and C++ do not provide any memory protection
 - Out of bounds array references
 - Invalid pointers
 - Abuses of `malloc` and `free`
- Can lead to nasty bugs
 - Depends on system and compiler
 - Action at a distance
 - Corrupted object logically unrelated to one being accessed
 - Effect of bug may be first observed long after it is generated
- How can I deal with this?
 - Program in Java, Ruby or ML
 - Understand what possible interactions may occur
 - Tools to detect referencing errors (e.g. `valgrind`)


Runs quickly

```
void copyij(int src[2048][2048], int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

Runs 21 times slower

```
void copyji(int src[2048][2048], int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

- 21 times slower (Pentium 4)
- Memory is organized hierarchically
- Performance depends on access patterns
 - Including stepping through multi-dimensional array



Reality #4

Performance is about more than asymptotic complexity

- Constant factors matter too!
- Even exact op count doesn't predict performance
 - Easily 10:1 performance range based on code written
 - Optimize at multiple levels: algorithm, data representations, procedures, loops
- Must understand system to optimize
 - How programs compiled and executed
 - How to measure performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Reality #5

Computers do more than run programs

- Need to get data in and out
 - I/O system critical to reliability and performance
- Communicate with over networks
 - Many system-level issues arise
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues